

# Intro to Kubernetes

John Zaccone

[john.zaccone@ibm.com](mailto:john.zaccone@ibm.com)

**IBM Developer**

# Lab

<https://github.com/jzaccone/kube101-1/tree/master/workshop/Lab1>

# Agenda

- What is Kubernetes?
- How was Kubernetes created?
- Where is Kubernetes?
- Kubernetes Architecture
- Resource Model
- Kubernetes in Action
- Kubernetes at IBM

# What is Kubernetes?

## **Container Orchestrator**

- Provision, manage, scale applications
- Manage infrastructure resources needed by applications
  - Volumes
  - Networks
  - Secrets
  - And many many many more...
- Declarative model
  - Provide the "desired state" and Kubernetes will make it happen
- What's in a name?
  - Kubernetes (K8s/Kube): "Helmsman" in ancient Greek

# How was Kubernetes created

- Based on Google's Borg & Omega
- Open Governance
  - Cloud Native Compute Foundation
- Adoption by Enterprise
  - RedHat, Microsoft, IBM and Amazon

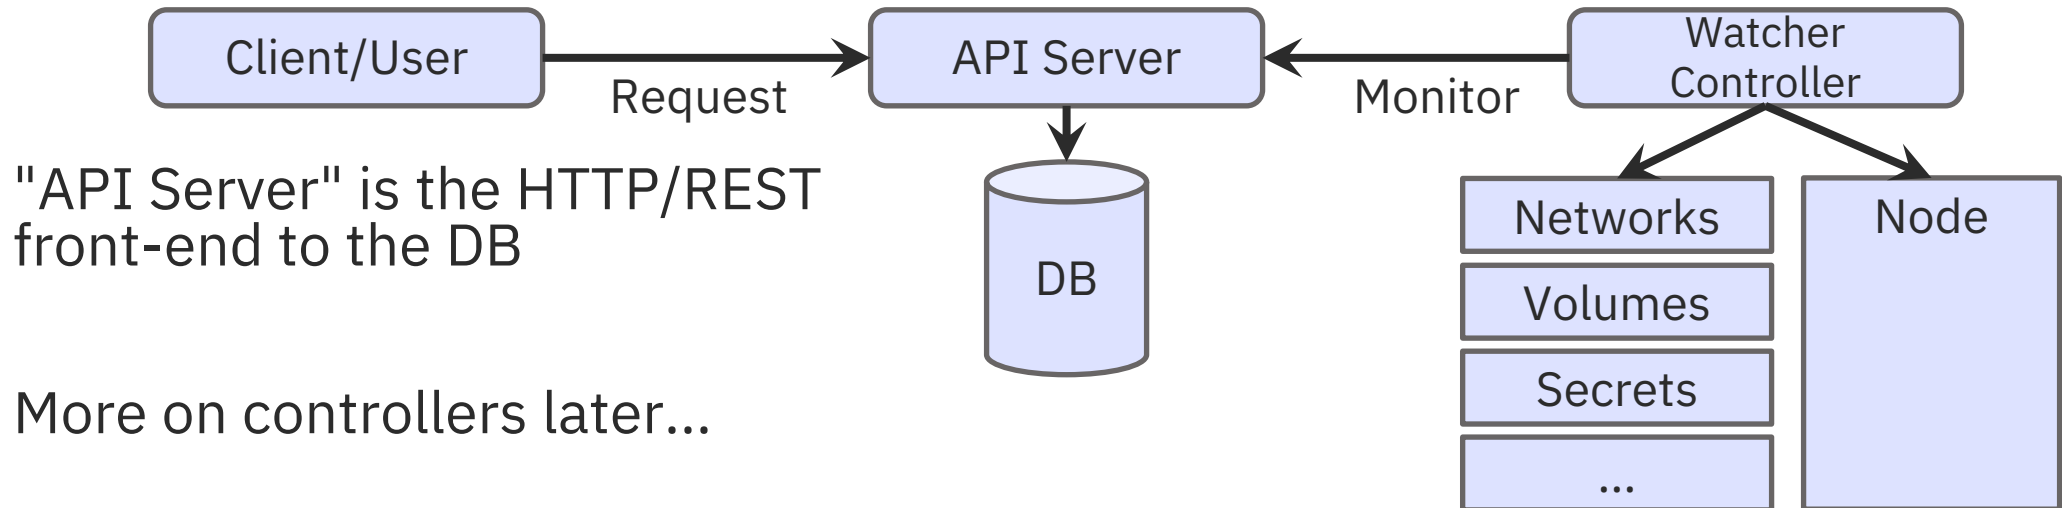
# Where is Kubernetes

Main Website - <http://kubernetes.io>

- Source Code - <https://github.com/kubernetes>
- [Youtube Channel](#)
- Many [SIG's](#)(Special Interest Groups), Zoom

# Kubernetes Architecture

- At its core, Kubernetes is a database (etcd).  
With "watchers" & "controllers" that react to changes in the DB.  
The controllers are what make it Kubernetes.  
This pluggability and extensibility is part of its "secret sauce".
- DB represents the user's desired state
  - Watchers attempt to make reality match the desired state



# Kubernetes Resource Model

## A resource for every purpose

- Config Maps
- Daemon Sets
- **Deployments**
- Events
- Endpoints
- Ingress
- Jobs
- Nodes
- Namespaces
- **Pods**
- Persistent Volumes
- Replica Sets
- Secrets
- Service Accounts
- **Services**
- Stateful Sets, and more...

- Kubernetes aims to have the building blocks on which you build a cloud native platform.
- Therefore, the internal resource model **is** the same as the end user resource model.

### Key Resources

- Pod: set of co-located containers
  - Smallest unit of deployment
  - Several types of resources to help manage them
  - Replica Sets, Deployments, Stateful Sets, ...
- Services
  - Define how to expose your app as a DNS entry
  - Query based selector to choose which pods apply



# Kubernetes Client

CLI tool to interact with Kubernetes cluster

Platform specific binary available to download

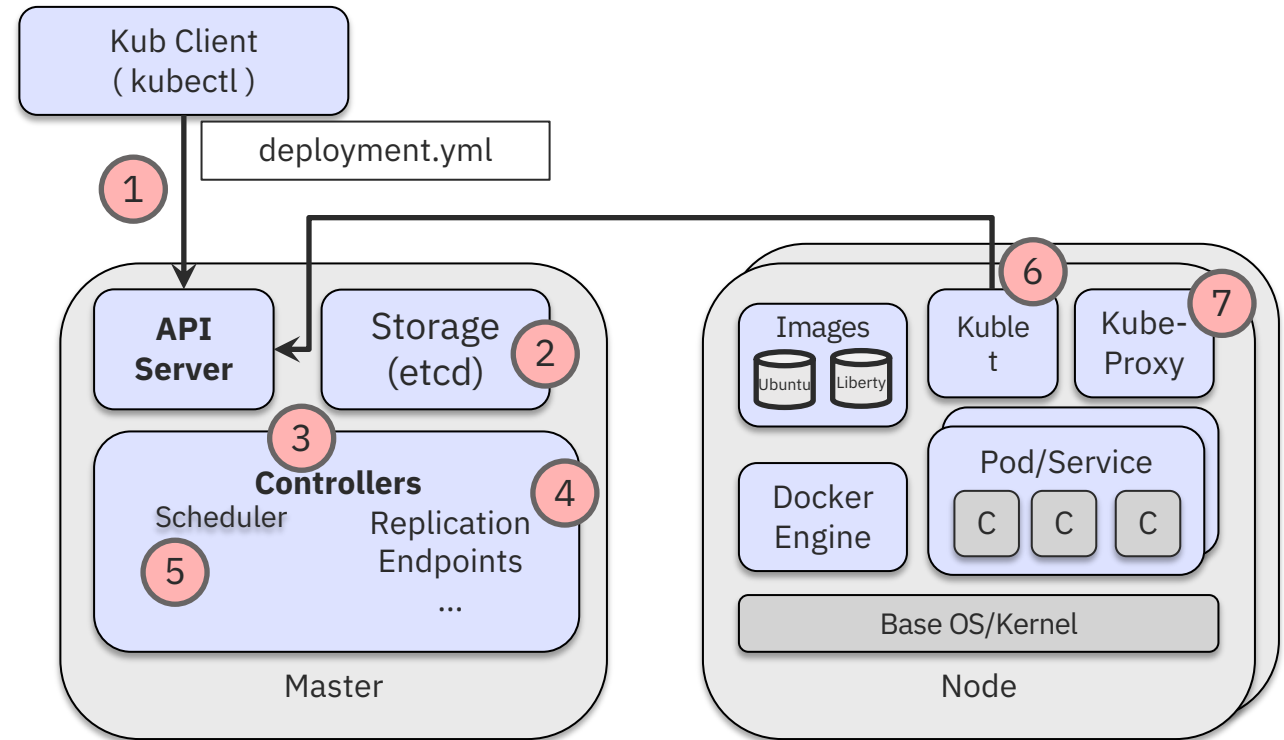
– <https://kubernetes.io/docs/tasks/tools/install-kubectl>

The user directly manipulates resources via json/yaml

```
$ kubectl (create|get|apply|delete) -f myResource.yaml
```

# Kubernetes in Action!

1. User via "kubectl" deploys a new application
2. API server receives the request and stores it in the DB (etcd)
3. Watchers/controllers detect the resource changes and act upon it
4. ReplicaSet watcher/controller detects the new app and creates new pods to match the desired # of instances
5. Scheduler assigns new pods to a kubelet
6. Kubelet detects pods and deploys them via the container runing (e.g. Docker)
7. Kubeproxy manages network traffic for the pods – including service discovery and load-balancing



# Kubernetes @ IBM

## Offerings / Plans

IBM Cloud Container Service – Docker containers orchestrated by K8s

ICP – IBM Cloud Private

IBM Cloud service teams use Kubernetes to host

## Key Development Activities

Service Catalog (co-lead)

Contributor Experience

Networking & Istio (co-lead)

ContainerD integration (co-lead)

Storage

Performance

# ICP installation

# Installing the IBM Cloud Private "developer edition"

Community effort started by DBG on [github.com/IBM](https://github.com/IBM) organization

Open terminal and clone this project:

```
$ git clone https://github.com/timroster/deploy-ibm-cloud-private.git
```

```
$ cd deploy-ibm-cloud-private
```

```
$ git checkout all_features
```

```
$ vagrant up
```

Installation will take 30-40 minutes