

# Microservices

## Principles, Patterns and Implementation considerations

**IBM Developer**


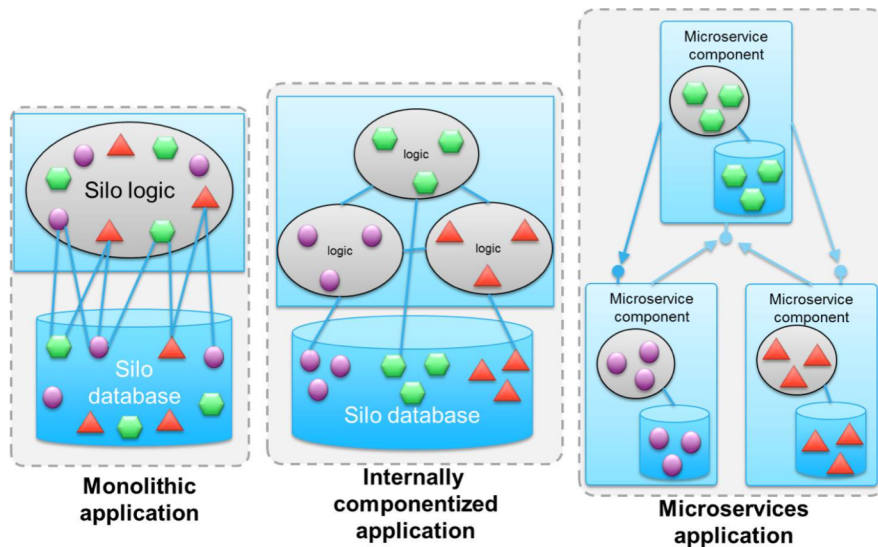
WW Developer Advocate Team

# Microservices and its Impact on CIO's Agenda

*"...the microservice architectural style... is an approach to developing a single application as a suite of small services, each running in its own process and communicating with lightweight mechanisms, often an HTTP resource API. ..."*

Martin Fowler

<http://martinfowler.com/articles/microservices.html>



## Agility and productivity

- The team that is developing the microservice can completely understand the codebase
- They can build, deploy, and test it independently of other components in much faster iteration cycles



## Scalability

- The Microservices development team can scale the component at run time independently of other microservice components
- These Microservices take advantage of the elastic capabilities of cloud-native environments that have cost-effective access to enormous resources



## Resilience

- The separate run time immediately provides resilience that is independent of failures in other components
- Technologies such as containers enable microservice components to fail quickly and independently, instead of taking down whole areas of unrelated functionality

# Small and Focused

- ✓ Business Oriented
- ✓ Microservices need to focus on a unit of work, and as such they are small. There are no rules on how small a microservice must be but making them too small can cause latency.
- ✓ A microservice needs to be treated like an application or a product. It should have its own source code management repository, and its own delivery pipeline for builds and deployment.
- ✓ Microservices powered applications provide light weight communication mechanism through the use of technology agnostic APIs.
- ✓ Microservices should have independent storage, should be independently changeable, should be independently deployable, should support distributed transactions.
- ✓ Reuse isn't the only business motivation for microservices. There are others, such as localized optimizations to improve user interface (UI) responsiveness and to be able to respond to customer needs more rapidly.

# Microservices – Fundamentals


## Microservices Architecture

Simplistically, microservices architecture is about breaking down large silo applications into more manageable fully decoupled pieces

### SOA done well



Monolithic application




Microservices application

A *microservice* is a granular decoupled component within a broader application

## Encapsulation is key


Related logic and data should remain together, and which means drawing strong boundaries between Microservices.

### Business Centric




## Microservice as API enabler

The “service” in “microservice” refers to the granularity of the **components**, not the exposed interfaces



An application split into Microservices may well expose the same APIs as its monolithic equivalent

# Polyglot Persistence



Ideally, you want every microservice to manage its own database. This enables polyglot persistence, using different databases (for example, Cloudant versus MongoDB, both of which are NoSQL) and different types of data stores (such as Structured Query Language (SQL), NoSQL, graph).

However, we might need to have more than one microservice use the same database for one of many reasons, for example, to preserve the ACID nature of a transaction that would otherwise be distributed across microservices and databases.

Database Type	Examples	When best used
Key-Value Databases	Redis by Compose, Bluemix Data Caching Service, Memcached, Amazon DynamoDB,	Storing interaction data, user preferences, simple shopping carts. Not great at set operations.
Document Databases	Cloudant, MongoDB by Compose	Event logging, content management, analytics. Not great at complex queries.
Column-Family Databases	Cassandra, Db2 BLU components for DB2 as a Service	Event logging, counters, blogs. Not compatible with ACID txns.
Graph Databases	Bluemix Graph DB, Neo4J, OrientDB,	Social Networks, Location-based Services. Not optimal for certain types of bulk updates.

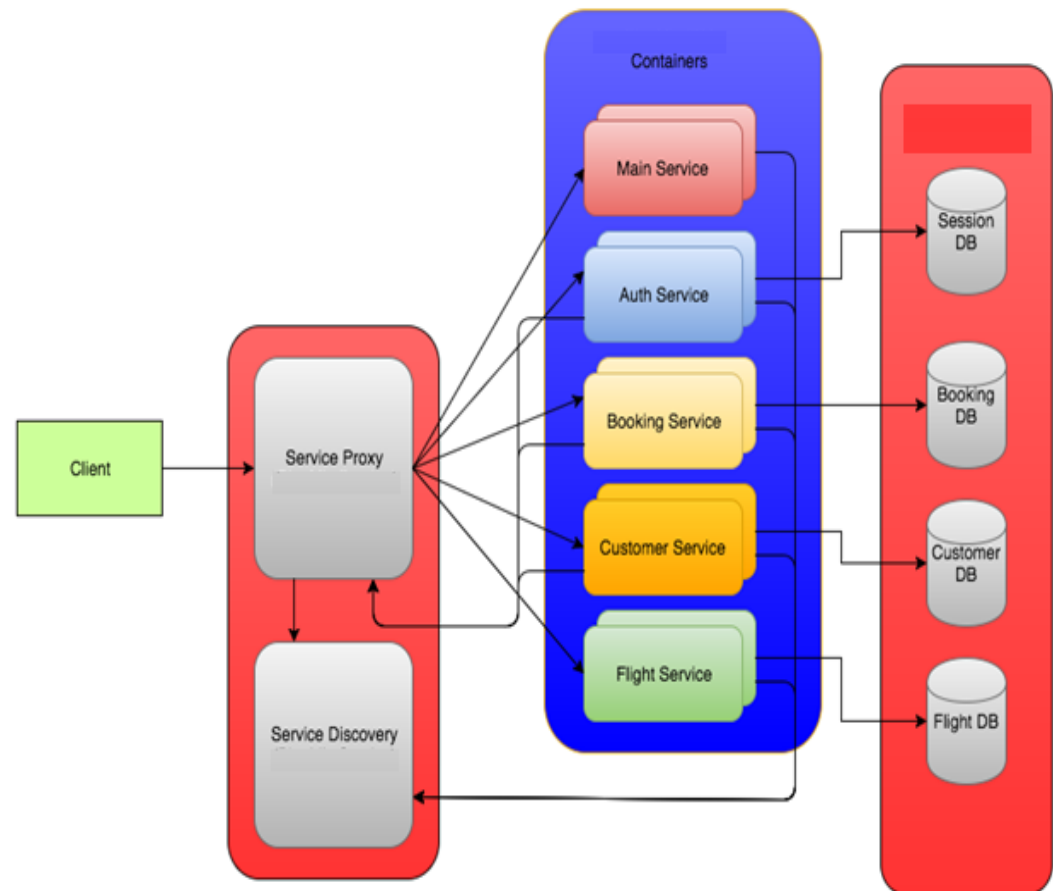
# Microservice Solution Overview

Here is a hypothetical Airline solution broken up into the following standalone services.

- Main app - presentation layer
- Auth Service - manages sessions (login, logout, validation)
- Customer Service - manages customer Info (get, update, validate password)
- Flight Service - manages flight data
- Booking Service - manages booking data
- Support Service - a simple dummy websocket chatroom (ma add a cognitive service here)

- Multiple instances of each service can be created.
- A service proxy is setup in front of the services to route/load balance them.
  - The Auth, Booking, and Customer services also need to know how to reach the proxy because they call other services.
  - Service Discovery manages the micro-services.

- The DBs can be separate instances (recommended) or one instance.



# Key tenets of a Microservices Architecture

## **Each microservice should be loosely coupled with other microservices**

- Large monoliths are broken down into many small services
  - Each service runs in its own process. The applicable cloud rule is one service per container
- Services are optimized for a single function
  - There is one and only one business function per service
- Communication via REST API and message brokers
  - Avoid tight coupling introduced by communication through a database.
- Per-service CI/CD
  - Services evolve at different rates.
  - You let the system evolve but set architectural principles to guide that evolution.
- Per-service HA/clustering decisions
  - Not one size fits all. Not all services need to scale; others require auto-scaling up to very large numbers.

## Microservices Design principles

High  
Cohesion

Loose  
Coupling

Business  
centric

Resilience

Observable

Automated

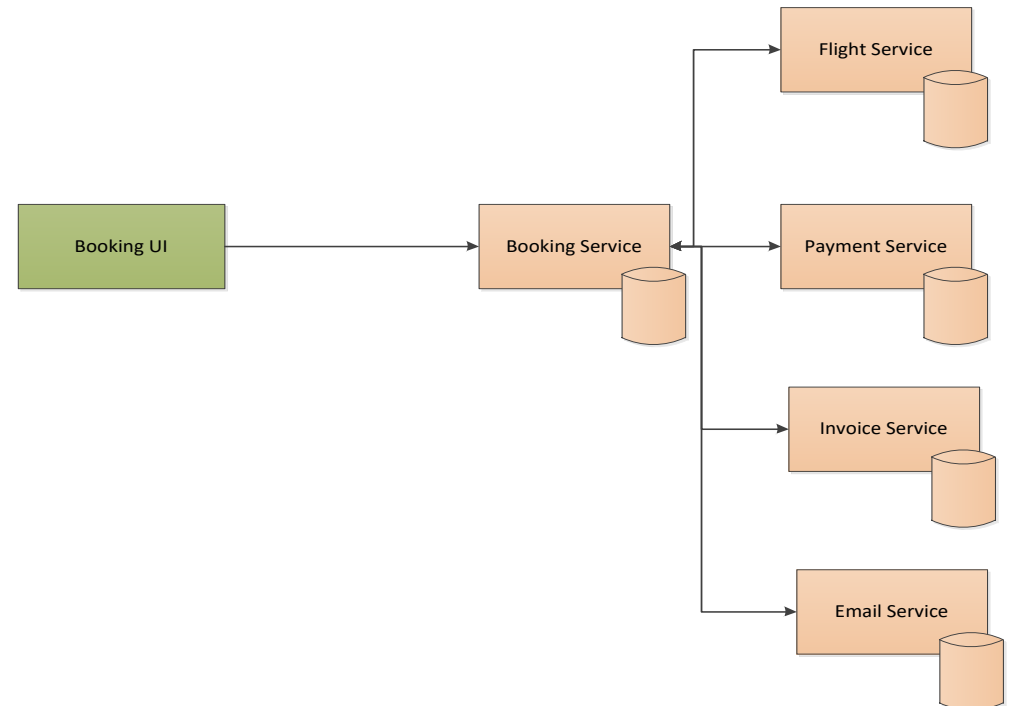


# High Cohesion

**Each microservices should have a single function or focus. example domain centric.**

Revisit our flight data example:

- Booking Service – The booking service allows you to browse flights (flight service), book tickets, payment service, generate invoice, send emails etc.
- In its simplest form, It could be implemented as seen in the first diagram where the booking UI and service are tightly coupled and hence “less” cohesive.
- Using the principles of high cohesion, we can further divide this functionality into multiple smaller, independent and focused micro services.
- High cohesion would allow us to manage each of these functionalities better, example if we change payment logic or email templates then it would not change the booking service and invoice service.



# Loose Coupling

## **Each microservice should be loosely coupled with other microservices**

Decouple microservices over the network:

- Synchronous Communication
- Asynchronous Communication

Use Open communication protocols

- HTTP/REST is standard being technology agnostic
- Avoid Client libraries resulting in RMI/EJB kind of communication

Standard interface for communication

- Have agreed interface for microservice communication
- Decouple internal representation of data within a microservice from how data is exchanged between two services.

Develop versioning strategy

- New changes should not break existing contract
- Should be backward compatible
- If breaking changes are unavoidable, then have parallel versions in play.
- Use semantic versioning to relay compatibility – Major, Minor, Patch
- Provide a migration strategy to consumers when new versions (not backward compatible) are available.

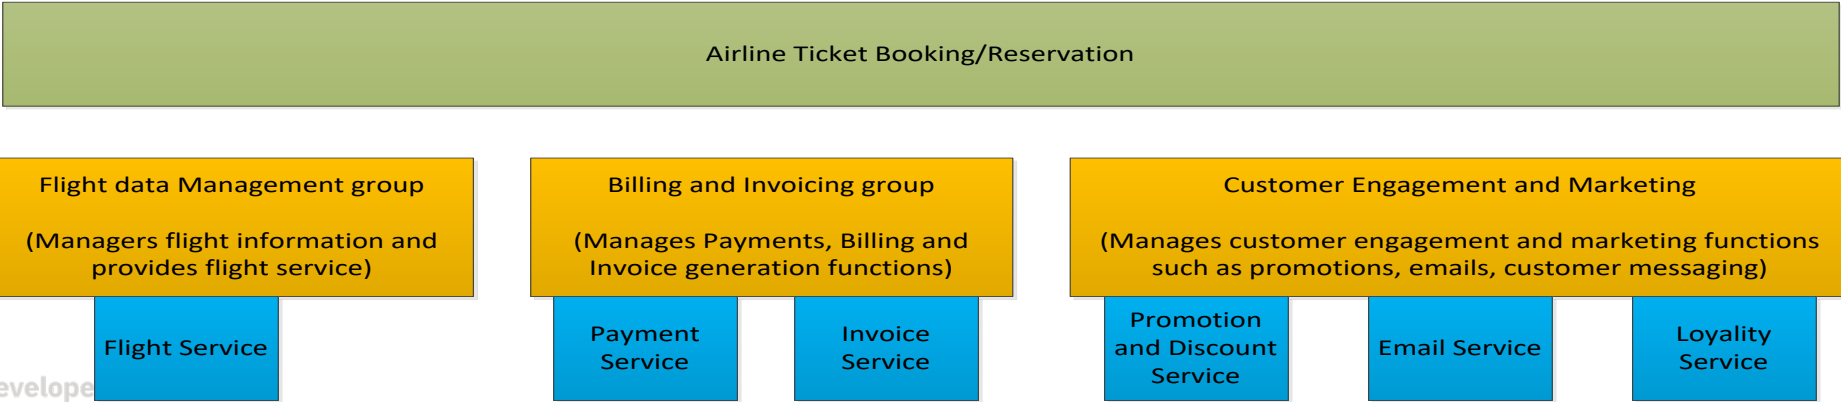

# Business Domain Centric

**Each microservices should map to a business domain or business function**

Microservices can be defined iteratively on the basis of business domain or business functions. These domains can initially be defined in a coarse manner and then further split into possible business functions.

Remember high cohesion and loose coupling principles when splitting microservices. This also helps us align and group services better with the right teams within a large organization.

Business services can delegate to other business services. Avoid circular dependencies. Careful that each service still implements a complete task. Not separate layers



# Resilience

## **Build microservices for resiliency**

Design for failures for down stream service.

When happens on failures

- Degrade functionality
- Default functionality

Design for fail fast

- Consider the usage/configuration of retries and timeouts. A long hanging transaction will not lead to a good user experience.
- Retries and Timeouts also help with network failures along with service failures

# Observable – Monitoring and Logging

## Centralized Monitoring

Use Real time monitoring

Monitor for operational indicator

- CPU, memory etc

Have services expose metrics which can be centrally logged and monitored

- Response times, Exceptions, Time-outs, Retries

Aggregate monitoring data

Use tools to help allow data aggregation, visualization to help see patterns and alert triggers.

## Centralized Logging

Critical in context of Microservices with multiple microservices having multiple instances.

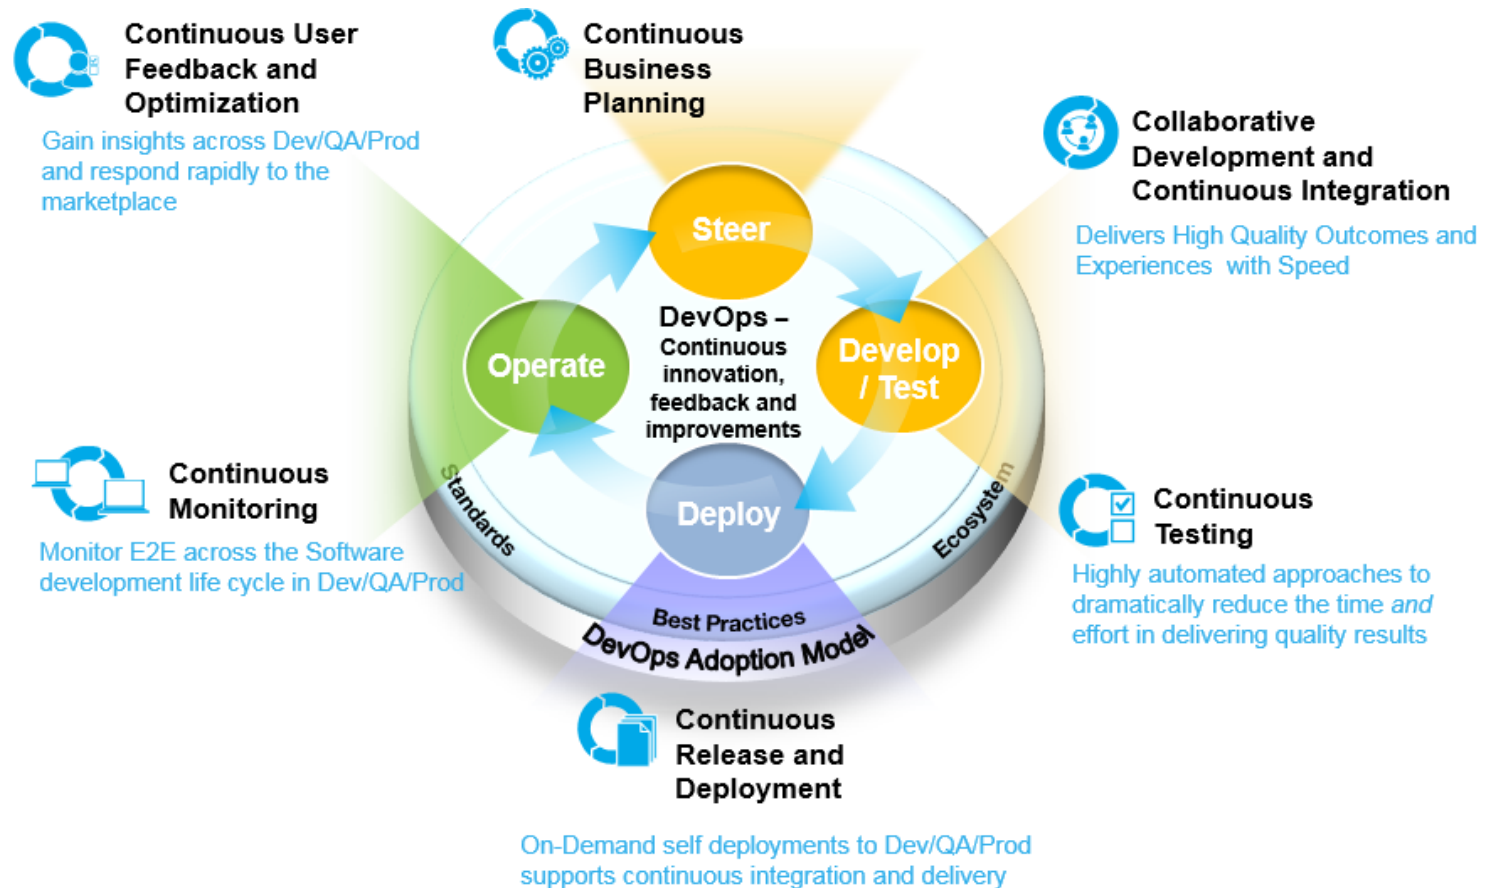
Critical to understand the lifecycle of how a request was processed in a distributed manner.

Include data such as host name, correlation id to help enable tracing a transaction across multiple services.

The log format should be consistent to help allow manage and query these logs later.

# Automation

## Devops principles for developing microservices



# Microservices Implementation Considerations

Communication

Hosting

Registration  
& Discovery

Monitoring

Logging

Caching

Scaling

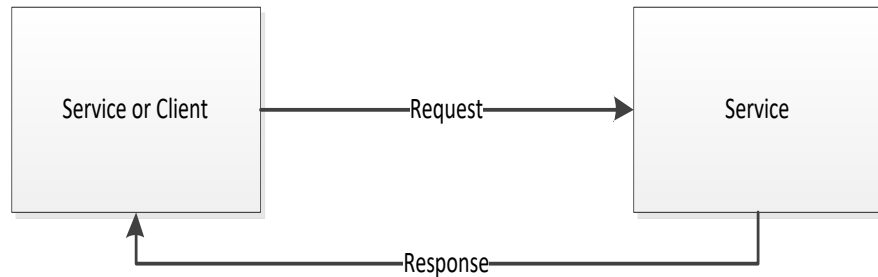
APIs

Automation

# Communication patterns


## Synchronous Communication

- HTTP/REST – Technology agnostic
- Have to wait until a response is received
- Design for failures – use timeouts or retries



## Asynchronous Communication

- Event based communication – instead of connecting directly to service, the service will publish an event which can be consumed by another service to be able to process the request.
- No need to wait for the response
- Use technologies such as Message Queuing





# Hosting considerations

Virtualization

Containerization


Self Hosting

A common practice is to run one microservice per container.

# Registration and Discovery

## Service Registry Database

- What to register?
  - Host, port, version
- When to register
  - On Startup
- When to deregister
  - On failures
- Cloud platforms makes it easy to manage registration and discovery. This has to be self managed mostly in a self hosted scenario.
- Service lookup patterns
  - Client connects to service registry – Client side discovery
  - Client connects to Gateway/Load Balancer to further connect to service registry – Server side discovery



# Monitoring and Logging Technology

## Centralized Monitoring tools

Example Nagios or New Relic

- Collect Metrics
- Aggregate
- Visualize data
- Send test transactions and monitor test transactions
- Monitor network
- Send alerts when required

## Centralized Logging

Example Logstash, Elastic log, splunk

The logs needs to be pushed to these tools which can then be stored in a database. Capabilities provided are:

- Push logs
- Query log
- Aggregate log from multiple servers
- Provide logging libraries for client side
- Visualize data
- Supports standardized logging

# Performance Considerations

## Scaling

- Horizontal Scaling is preferred and readily supported through cloud capabilities.
- Scaling can be automated or made on-demand
- Automated

## Caching

- Helps improve performance
- Ideally implemented at API gateway layer
- Can also be done on client layer example for SPA applications
- Can also be done at service layer example for static data

## API Gateway

- Load balancing and Caching
- Centralized Security
- Centralized Analytics
- Centralized service discovery
- Centralized entry points

# References

<https://martinfowler.com/articles/microservices.html>

<http://www.redbooks.ibm.com/abstracts/sg248275.html?Open>

<https://www.nginx.com/blog/microservices-at-netflix-architectural-best-practices/>

<https://w3-03.sso.ibm.com/services/lighthouse/documents/D023183V03153V20>

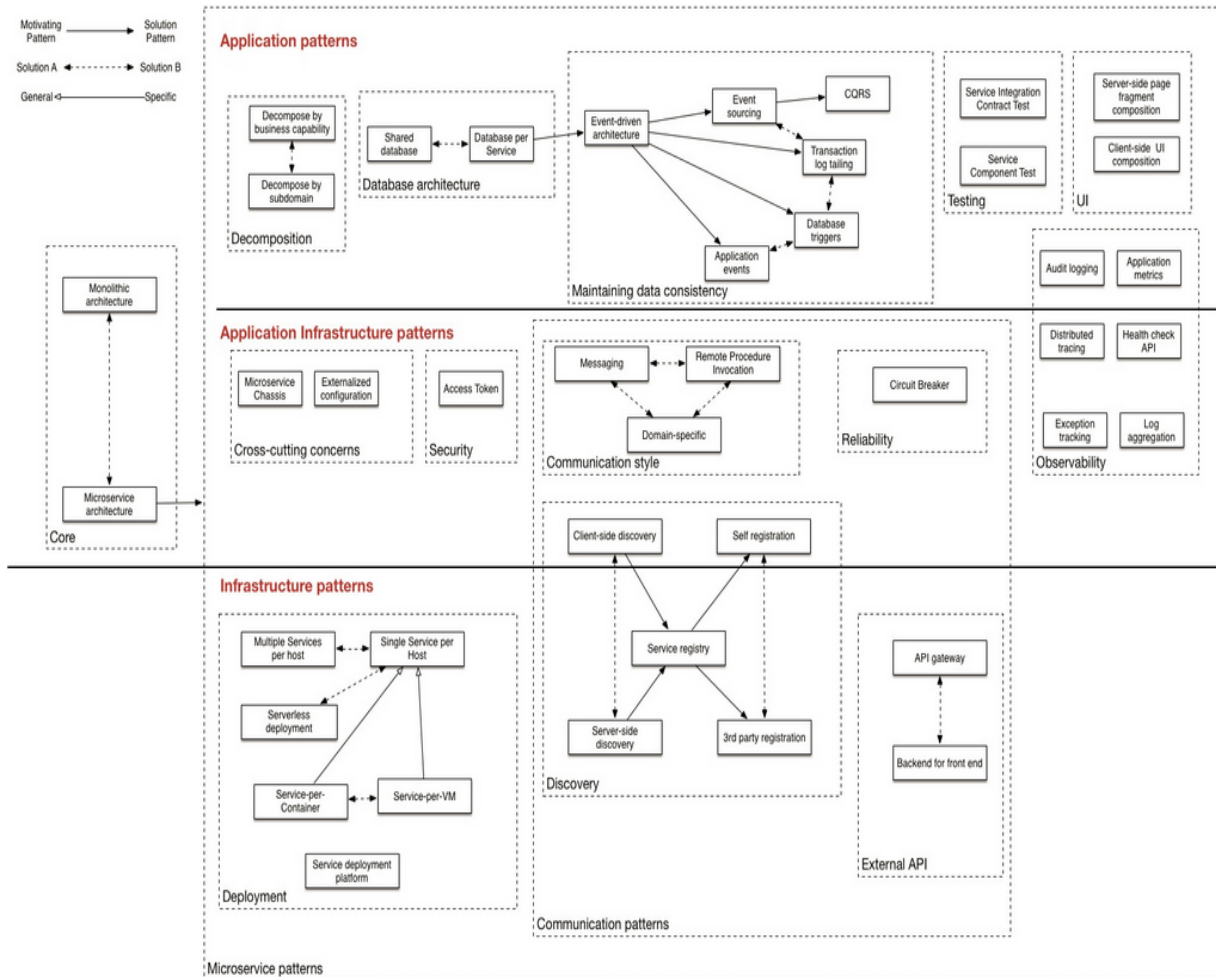
<https://12factor.net/>

<http://www.grahamlea.com/2016/08/distributed-transactions-microservices-icebergs/>

<http://microservices.io/patterns/index.html>

# Microservices Patterns

Microservices.io defines these patterns well - <http://microservices.io/patterns/index.html>



Service discovery  
[Client-side discovery](#)  
[Server-side discovery](#)  
[Service registry](#)  
[Self registration](#)  
[3rd party registration](#)

## Core patterns

- [Monolithic architecture](#)
- [Microservice architecture](#)

## Decomposition

- [Decompose by business capability](#)
- [Decompose by subdomain](#)

## Deployment patterns

- [Multiple service instances per host](#)
- [Service instance per host](#)
- [Service instance per VM](#)
- [Service instance per Container](#)
- [Serverless deployment](#)
- [Service deployment platform](#)

## Cross cutting concerns

- [Microservice chassis](#)
- [Externalized configuration](#)

## Communication style

- [Remote Procedure Invocation](#)
- [Messaging](#)
- [Domain-specific protocol](#)

## External API

- [API gateway](#)
- [Backend for front-end](#)

## Security

- [Access Token](#)

## Data management

- [Database per Service](#)
- [Shared database](#)
- [Event-driven architecture](#)
- [Event sourcing](#)
- [Transaction log tiling](#)
- [Database triggers](#)
- [Application events](#)
- [CQRS](#)

## Observability

- [Log aggregation](#)
- [Application metrics](#)
- [Audit logging](#)
- [Distributed tracing](#)
- [Exception tracking](#)
- [Health check API](#)

## UI patterns

- [Server-side page fragment composition](#)
- [Client-side UI composition](#)

## Testing

- [Service Component Test](#)
- [Service Integration Contract Test](#)

## Reliability

- [Circuit Breaker](#)

